



BMB5113

COMPUTER VISION


IMAGE PROCESSING I

Digital Images in General

- Image is a function of four variables $f(x,y; t, \lambda)$
 - For color image, λ takes three different values corresponding to red, green and blue components
 - t is a time variable for a sequence of frames
 - for constant λ (black and white) and a constant t
 - f becomes a function of two spatial variables
 - **Binary Image :**
 - If there are just two values, e.g. black and white, we usually represent them by 0 and 1.
 - **Gray-scale Image :**
 - If there is just a single plane of color, usually represented in between 0 and 255 (for int/uint) or 0 and 1 (for float or double)

A Digital Image

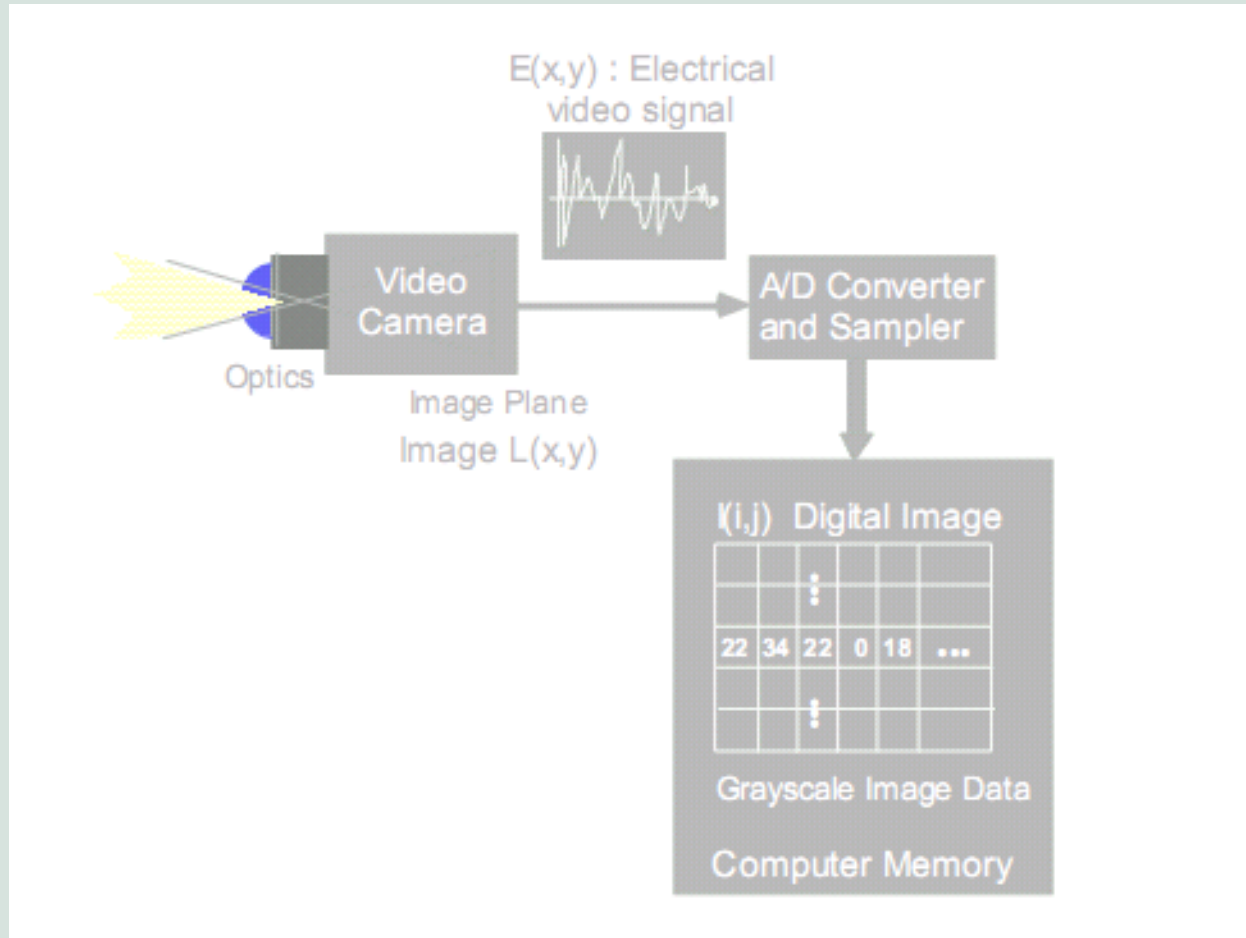
- 2-D function of intensities
 - can be represented as $f(x,y)$
 - $(x,y) \rightarrow$ 2-D coordinate pair
 - $f(x,y) \rightarrow$ the image intensity value at (x,y)



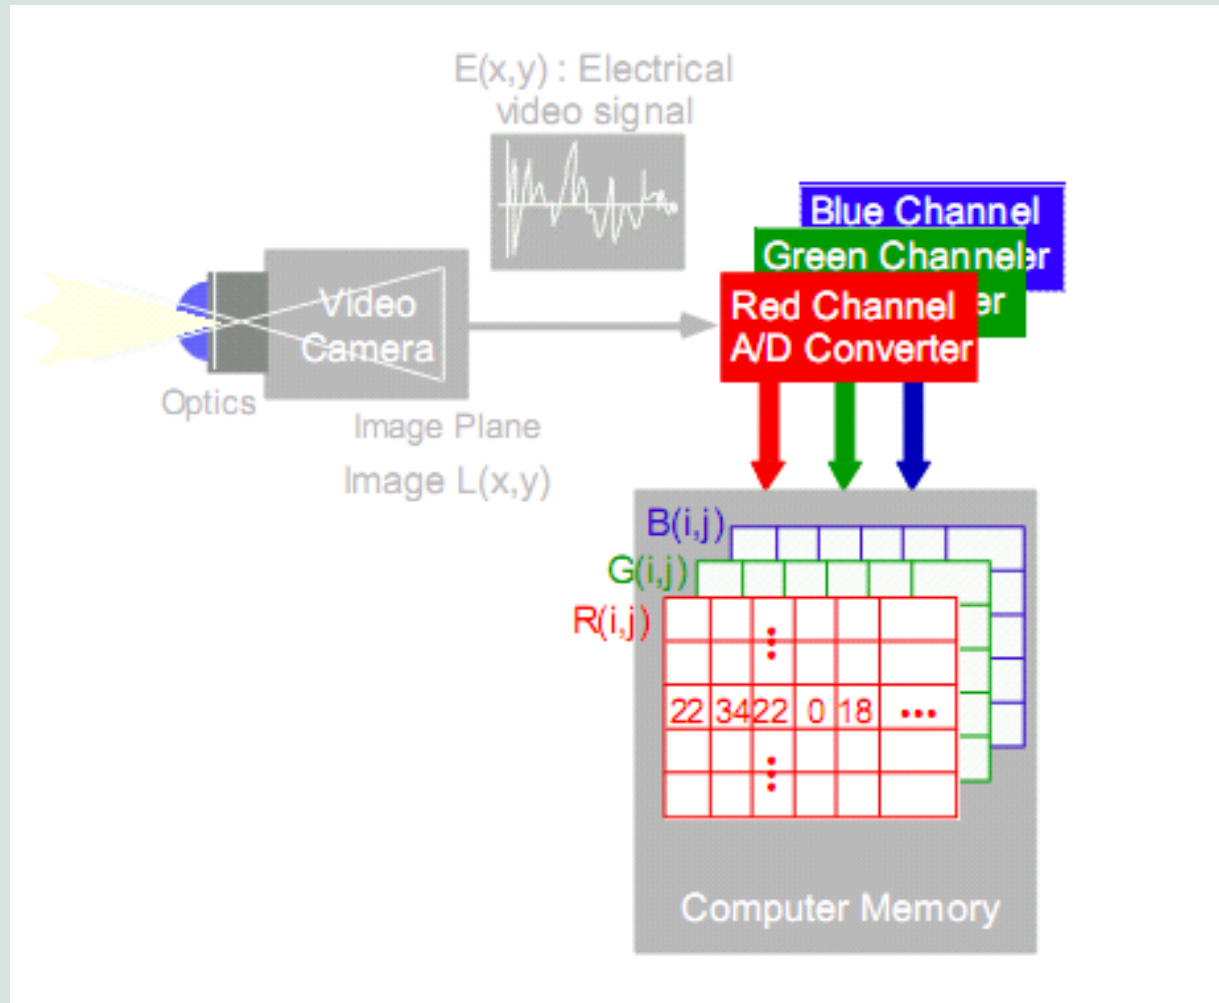
Digitization of An Image

- Filtering
 - For successful reconstruction $f_s \geq 2f_{\max}$
- Sampling
 - Taking samples at specific intervals from 2-D spatial plane
- Quantization
 - Adjustment for intensities to have specific values
- Digitization
 - Converting the intensity values to digital form

Grayscale Image Sensing Systems:

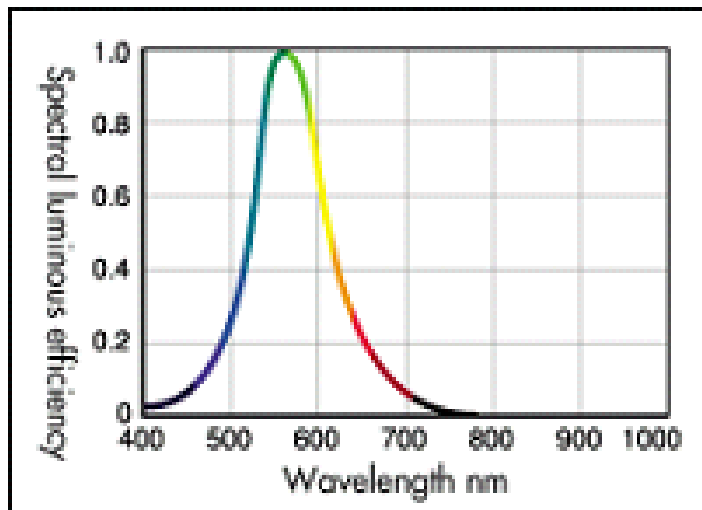


Color Image Image Sensing Systems:



CCD cameras are much more sensitive than the eye

Human Eye



CCD Camera


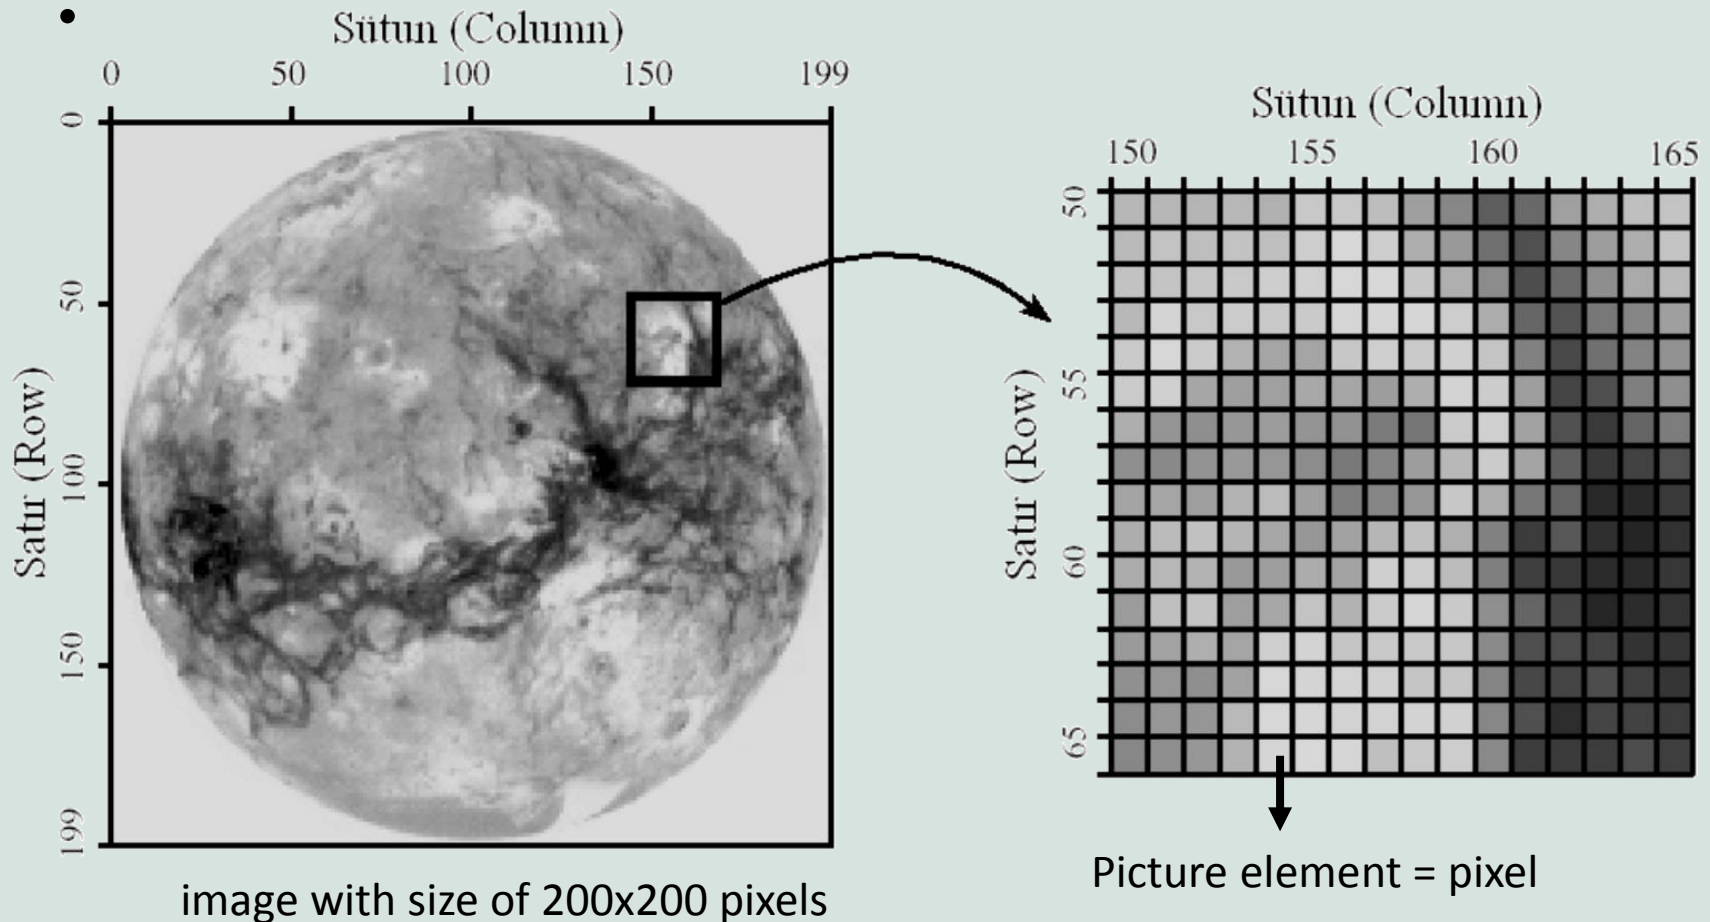


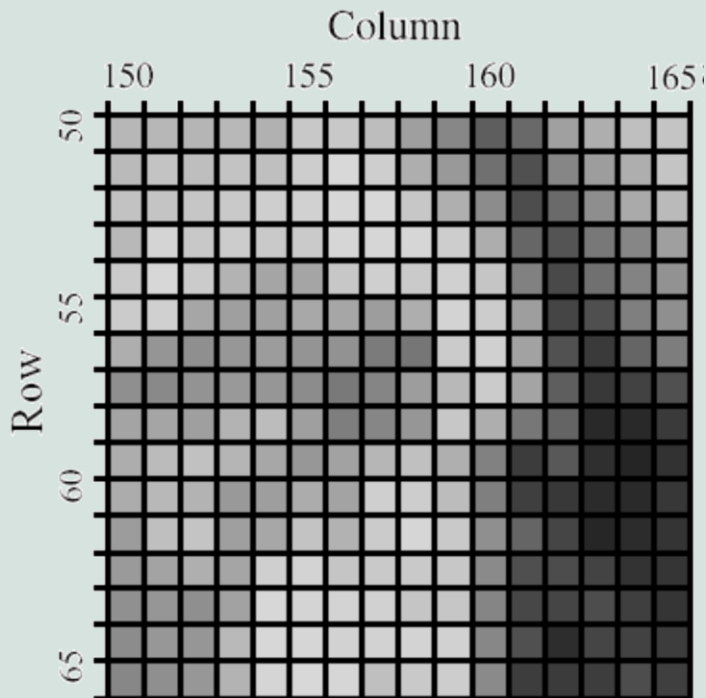
Image Pixels

- After sampling and quantization, a digital image is a rectangular array of integer values.



Pixel Intensity Values

- Intensity values are
 - non-negative
 - called gray levels



		Column															
		150				155				160				165			
Row	50	183	183	181	184	177	200	200	189	159	135	94	105	160	174	191	196
		186	195	190	195	191	205	216	206	174	153	112	80	134	157	174	196
		194	196	198	201	206	209	215	216	199	175	140	77	106	142	170	186
		184	212	200	204	201	202	214	214	214	205	173	102	84	120	134	159
		202	215	203	179	165	165	199	207	202	208	197	129	73	112	131	146
	55	203	208	166	159	160	168	166	157	174	211	204	158	69	79	127	143
		174	149	143	151	156	148	146	123	118	203	208	162	81	58	101	125
		143	137	147	153	150	140	121	133	157	184	203	164	94	56	66	80
		164	165	159	179	188	159	126	134	150	199	174	119	100	41	41	58
		173	187	193	181	167	151	162	182	192	175	129	60	88	47	37	50
	60	172	184	179	153	158	172	163	207	205	188	127	63	56	43	42	55
		156	191	196	159	167	195	178	203	214	201	143	101	69	38	44	52
		154	163	175	165	207	211	197	201	201	199	138	79	76	67	51	53
		144	150	143	162	215	212	211	209	197	198	133	71	69	77	63	53
		140	151	150	185	215	214	210	210	211	209	135	80	45	69	66	60
	65	135	143	151	179	213	216	214	191	201	205	138	61	59	61	77	63

Picture element = pixel

Spatial Resolution

256x256 pixel

128x128 pixel

64x64 pixel

72 dpi



36 dpi



18 dpi



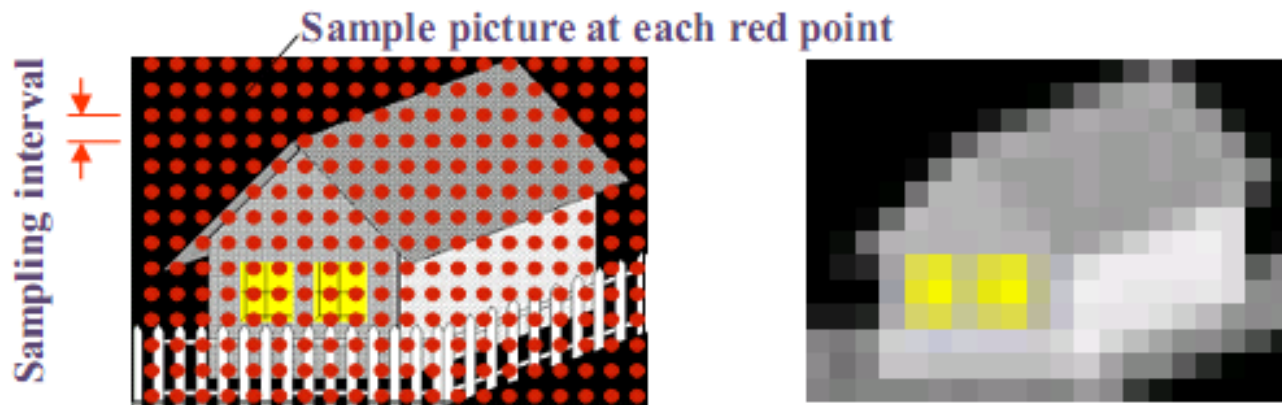
32x32 pixel

9 dpi




- Images with same physical dimensions (3.6''x3.6'')
- But with different resolutions
- Resolution = $\frac{\text{\# of pixels}}{\text{unit physical dimension}}$
- Resolution units
 - dpi: dots per inch
 - ppi: pixels per inch

Sampling (Resolution)



Coarse Sampling: 20 points per row by 14 rows



Finer Sampling: 100 points per row by 68 rows

Bit depth of an Image

L = 8 bit



L = 7 bit



L = 6 bit



L = 5 bit



L = 4 bit



L = 3 bit




L = 2 bit




L = 1 bit

Grayscale Quantization Level:



Color Image Quantization Level



$K=2$ (each color)



$K=4$ (each color)

VSCode Setup for Building Programs in bm5113 Environment

- After environment setup arrange VSCode json files
 - "terminal.integrated.defaultProfile.windows": "Command Prompt",
 - "python.condaPath": "C:\\ProgramData\\Anaconda3\\Scripts\\conda.exe"
- Run the following update
 - conda update anaconda-navigator
 - conda update navigator-updater

Python Image Functions: Read

- Reading an image:

```
from PIL import Image  
import numpy as np  
im = np.array(Image.open('./jpg/car1.jpg'))
```
- Converting image to grayscale after reading

```
pil_im = Image.open('empire.jpg').convert('L')
```
- Row and column dimensions of an image:

```
print(im.shape)  
[M,N,C] = im.shape
```


Python Image Functions: Display

- Displaying an image:

```
import matplotlib.pyplot as plt
```

```
imgplot = plt.imshow(im, cmap='gray')
```

- cmap: the colormap used to display image
- by default 'viridis'
- vmin=0, vmax=255 parameters can be used when displaying
- Interactive displaying of the pixel values possible

Python Image Functions: Write

- Writing an image:

```
pil_img = Image.fromarray(im)  
pil_img.save('./jpg/car1_copy.jpg')
```

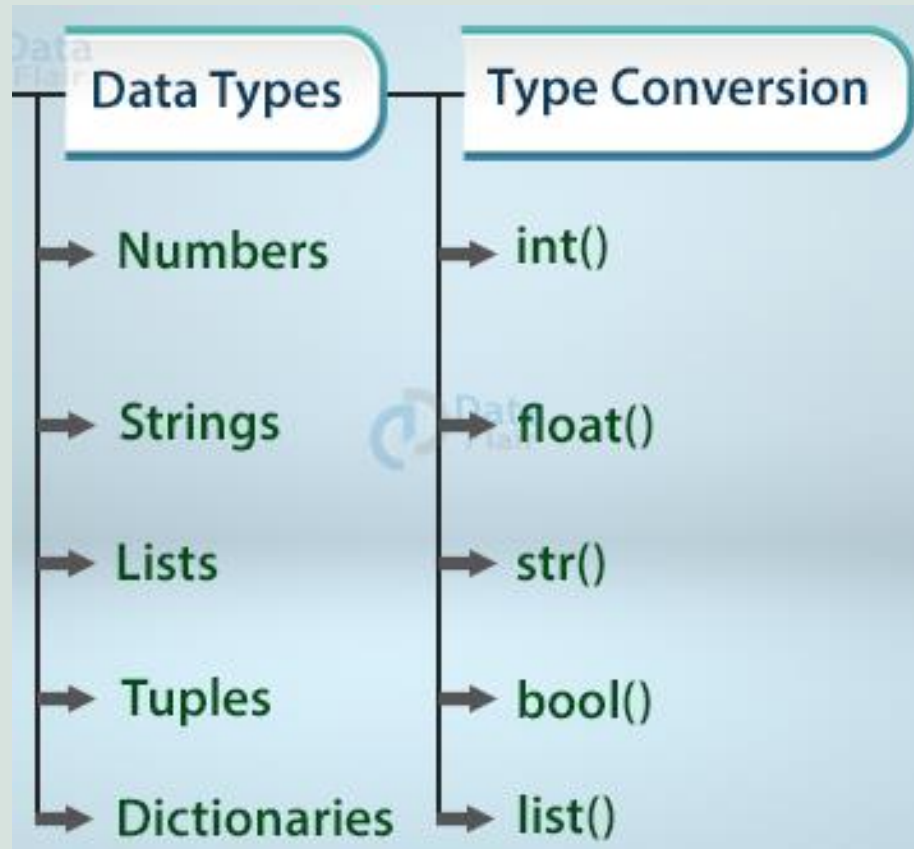
- Minimum and maximum values of the image

```
print(im_f[:, :, :].min())  
print(im_f[:, :, :].max())
```

Python Variable Types

Class	Description	Immutable?
bool	Boolean value	✓
int	integer (arbitrary magnitude)	✓
float	floating-point number	✓
list	mutable sequence of objects	
tuple	immutable sequence of objects	✓
str	character string	✓
set	unordered set of distinct objects	
frozenset	immutable form of set class	✓
dict	associative mapping (aka dictionary)	

Python Conversion of Types



Python Arithmetic Operators

Operator	Description	Example
+ Addition	Adds values on either side of the operator.	$a + b = 30$
- Subtraction	Subtracts right hand operand from left hand operand.	$a - b = -10$
* Multiplication	Multiplies values on either side of the operator	$a * b = 200$
/ Division	Divides left hand operand by right hand operand	$b / a = 2$
% Modulus	Divides left hand operand by right hand operand and returns remainder	$b \% a = 0$
** Exponent	Performs exponential (power) calculation on operators	$a ** b = 10 \text{ to the power } 20$
//	Floor Division - The division of operands where the result is the quotient in which the digits after the decimal point are removed. But if one of the operands is negative, the result is floored, i.e., rounded away from zero (towards negative infinity) –	$9 // 2 = 4$ and $9.0 // 2.0 = 4.0$, $-11 // 3 = -4$, $-11.0 // 3 = -4.0$

Python Comparison Operators

- When
 - a=3
 - b=5

Operator	Description	Example
==	If the values of two operands are equal, then the condition becomes true.	(a == b) is not true.
!=	If values of two operands are not equal, then condition becomes true.	(a != b) is true.
<>	If values of two operands are not equal, then condition becomes true.	(a <> b) is true. This is similar to != operator.
>	If the value of left operand is greater than the value of right operand, then condition becomes true.	(a > b) is not true.
<	If the value of left operand is less than the value of right operand, then condition becomes true.	(a < b) is true.
>=	If the value of left operand is greater than or equal to the value of right operand, then condition becomes true.	(a >= b) is not true.
<=	If the value of left operand is less than or equal to the value of right operand, then condition becomes true.	(a <= b) is true.

Python Assignment Operators

Operator	Description	Example
=	Assigns values from right side operands to left side operand	<code>c = a + b</code> assigns value of <code>a + b</code> into <code>c</code>
<code>+=</code> Add AND	It adds right operand to the left operand and assign the result to left operand	<code>c += a</code> is equivalent to <code>c = c + a</code>
<code>-=</code> Subtract AND	It subtracts right operand from the left operand and assign the result to left operand	<code>c -= a</code> is equivalent to <code>c = c - a</code>
<code>*=</code> Multiply AND	It multiplies right operand with the left operand and assign the result to left operand	<code>c *= a</code> is equivalent to <code>c = c * a</code>
<code>/=</code> Divide AND	It divides left operand with the right operand and assign the result to left operand	<code>c /= a</code> is equivalent to <code>c = c / a</code>
<code>%=</code> Modulus AND	It takes modulus using two operands and assign the result to left operand	<code>c %= a</code> is equivalent to <code>c = c % a</code>
<code>**=</code> Exponent AND	Performs exponential (power) calculation on operators and assign value to the left operand	<code>c **= a</code> is equivalent to <code>c = c ** a</code>
<code>//=</code> Floor Division	It performs floor division on operators and assign value to the left operand	<code>c //= a</code> is equivalent to <code>c = c // a</code>

Python Bitwise Operators

- When $a=0011\ 1100$ $b=0000\ 1101$

Operator	Description	Example
& Binary AND	Operator copies a bit to the result if it exists in both operands	$(a \& b)$ (means $0000\ 1100$)
Binary OR	It copies a bit if it exists in either operand.	$(a b) = 61$ (means $0011\ 1101$)
^ Binary XOR	It copies the bit if it is set in one operand but not both.	$(a \wedge b) = 49$ (means $0011\ 0001$)
~ Binary Ones Complement	It is unary and has the effect of 'flipping' bits.	$(\sim a) = -61$ (means $1100\ 0011$ in 2's complement form due to a signed binary number.
<< Binary Left Shift	The left operands value is moved left by the number of bits specified by the right operand.	$a \ll 2 = 240$ (means $1111\ 0000$)
>> Binary Right Shift	The left operands value is moved right by the number of bits specified by the right operand.	$a \gg 2 = 15$ (means $0000\ 1111$)

Python Logical Operators

- When a=1 and b=1

Operator	Description	Example
and Logical AND	If both the operands are true then condition becomes true.	(a and b) is true.
or Logical OR	If any of the two operands are non-zero then condition becomes true.	(a or b) is true.
not Logical NOT	Used to reverse the logical state of its operand.	Not(a and b) is false.

Python Membership and Identity Operators

Operator	Description	Example
in	Evaluates to true if it finds a variable in the specified sequence and false otherwise.	x in y, here in results in a 1 if x is a member of sequence y.
not in	Evaluates to true if it does not finds a variable in the specified sequence and false otherwise.	x not in y, here not in results in a 1 if x is not a member of sequence y.

Operator	Description	Example
is	Evaluates to true if the variables on either side of the operator point to the same object and false otherwise.	x is y, here is results in 1 if id(x) equals id(y).
is not	Evaluates to false if the variables on either side of the operator point to the same object and true otherwise.	x is not y, here is not results in 1 if id(x) is not equal to id(y).

Python Image Operator Functions

- `from PIL import Image, ImageMath`
- `abs(image)`
 - Absolute value.
- `convert(image, mode)`
 - Convert image to the given mode.
 - The mode must be given as a string constant.
- `float(image)`
 - Convert image to 32-bit floating point.
 - This is equivalent to `convert(image, "F")`.
- `int(image)`
 - Convert image to 32-bit integer. This is equivalent to `convert(image, "I")`.
 - Note that 1-bit and 8-bit images are automatically converted to 32-bit integers if necessary to get a correct result.
- `max(image1, image2)`
 - Maximum value.
- `min(image1, image2)`
 - Minimum value.

Python

Generating and Processing an Image

```
img = np.zeros([100,100,3],dtype=np.uint8)
# generate a new black image
img = Image.new( 'RGB', (2000,2000), "black")
# generate the pixel map
pixels = img.load()
for i in range(img.size[0]):
    # for every pixel:
    for j in range(img.size[1]):
        #do some stuff that requires i and j as parameter
        if pixels[i,j] == (255, 0, 0):
            pixels[i,j] = (0, 0 ,0)
```

Python Special Variables

- `import math`
`math.nan`
`math.pi`
- `import sys`
`sys.float_info.epsilon`

Image Enhancement



Dark image



Image after enhancement

Pixelwise Operations: Illumination Adjustment

- General form: $g(x,y)=T[f(x,y)]$
- Add to or subtract from image intensities $f(x,y)$ a constant number b

$$g(x,y) = f(x,y) + b$$

Original

$B=-50$

$B=+50$



Pixelwise Operations: Contrast Adjustment

- Multiply image intensities $f(x,y)$ a constant number a

$$g(x,y) = af(x,y)$$

Original



$a < 1$




$a > 1$




Pixelwise Operations: Illumination & Contrast Adjustment

$$g(x,y) = af(x,y) + b$$





Pixelwise Operations: Simple Thresholding

$$g(x,y) = f(x,y) \geq t$$



Pixelwise Operations: Logarithmic and Gamma Transformations

- Logarithmic transformations $g = c * \log(1 + f(x, y))$
- Gamma corrections $g = c * f(x, y)^\gamma$



Python Adjusting Image Intensities

```
def imadjust(x,a,b,c,d,gamma=1):
```

```
    # Similar to imadjust in MATLAB.
```

```
    # Converts an image range from [a,b] to [c,d].
```

```
    # The Equation of a line can be used for this transformation:
```

```
    #  $y = ((d-c)/(b-a)) * (x-a) + c$ 
```

```
    # However, it is better to use a more generalized equation:
```


```
    #  $y = ((x-a)/(b-a))^{\text{gamma}} * (d-c) + c$ 
```

```
    # If gamma is equal to 1, then the line equation is used.
```

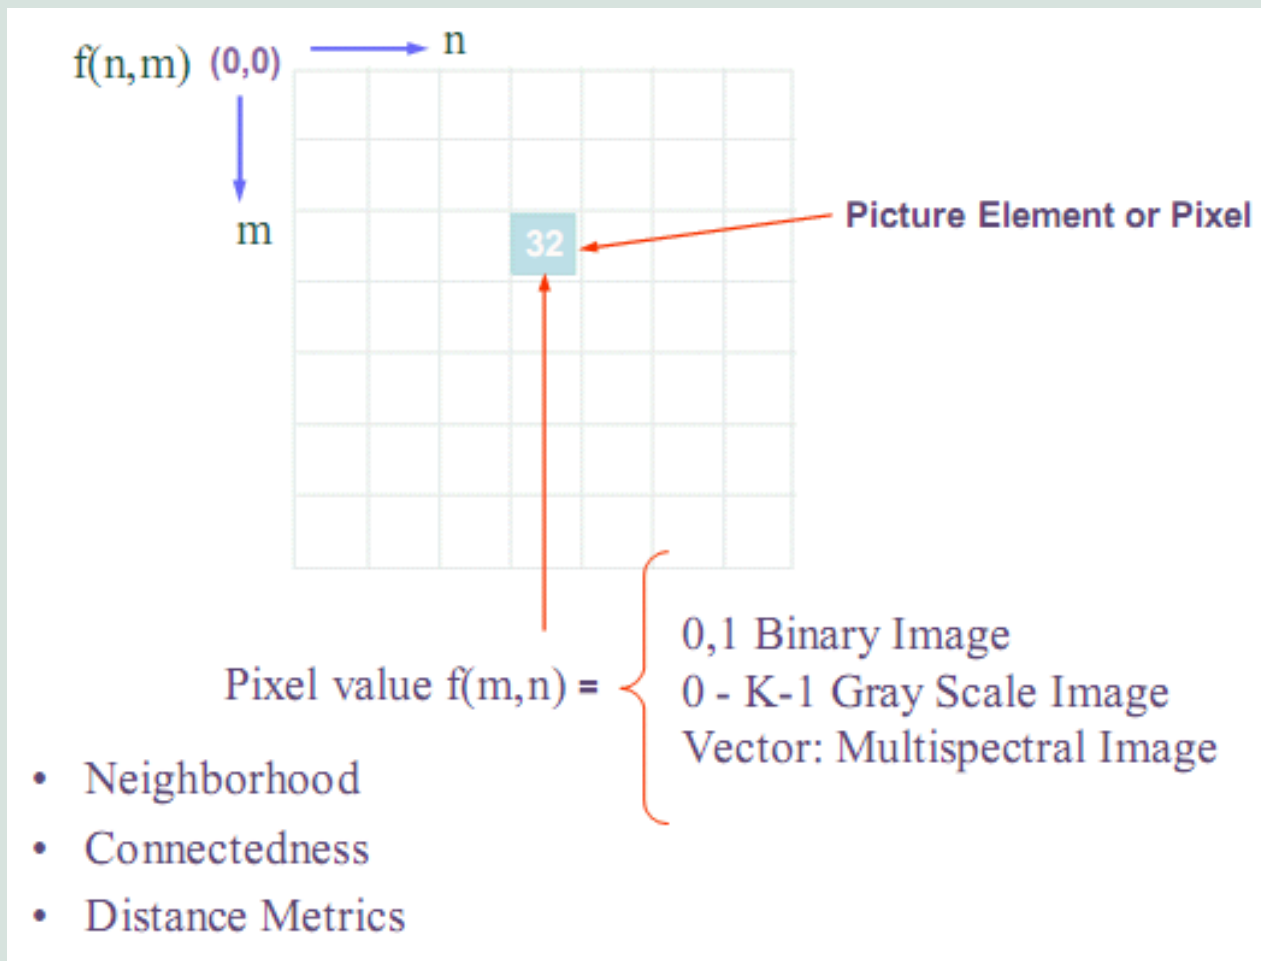
```
    # When gamma is not equal to 1, then the transformation is not  
    linear.
```

```
     $y = (((x - a) / (b - a)) ** \text{gamma}) * (d - c) + c$ 
```

```
    return y
```



Pixel-Related Concepts




Pixel Neighborhood

- Except on borders of the array, any point (m,n) has 8 neighbor pixels


$$\begin{bmatrix} (m-1, n-1) & (m-1, n) & (m-1, n+1) \\ (m, n-1) & (m, n) & (m, n+1) \\ (m+1, n-1) & (m+1, n) & (m+1, n+1) \end{bmatrix}$$

- Note that diagonal neighbors $\sqrt{2}$ units away from (m,n) while horizontal and vertical neighbors are only 1 unit away.

Pixel Connectedness




Four Neighbor




Eight Neighbor

Distance Metrics




Euclidean Distance

$$= \sqrt{(i-n)^2 + (j-m)^2}$$



City Block Distance


$$= |i-n| + |j-m|$$



Chessboard Distance

$$= \max[|i-n|, |j-m|]$$


Noise Removal: Red Blood Cell Image



- Many blood cells are separate objects
- Many touch – bad!
- Salt and pepper noise from thresholding

Removing Salt-and-Pepper Noise

- Change pixels all of whose neighbors are different:
 - see hole filled at bottom
- Delete objects that are tiny relative to targets:
 - see some islands removed at right



Neighborhood (Spatial) Operations for Enhancement

- Linear Filtering:
 - low-pass filter masks
- Nonlinear filtering:
 - median, max, min, and other order-statistic operators
- Morphological filters

Linear filtering

- Convolution with a filter mask is a weighted sum of pixel intensities

5	5	5	5	5	5	5
5	5	5	5	5	5	5
5	5	5	5	5	5	5
0	0	0	0	5	5	5
0	5	0	0	0	5	0
5	0	0	5	0	0	0
0	0	0	0	0	0	0

Input

$1/5 \times$

0	1	0
1	1	1
0	1	0

Mask
(Kernel)

Mean Filter

5	5	5	5	5	5	5
5	5	5	5	5	5	5
4	4	4	4	5	5	5
1	2	1	2	3	5	4
2	1	1	1	2	2	2
2	2	1	1	1	1	0
1	0	0	1	0	0	0


Output

Linear Filtering: Lowpass Filters

- Lowpass filters reduce noise,...
- But smooths edges
- Mean filter is a LPF
- Gaussian Filter is a better LPF

LP Filters: Gaussian

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Gaussian Filters are separable: you can filter rows then columns...

.006	.061	.242	.383	.242	.061	.006
------	------	------	------	------	------	------

or, apply 2D convolution:

$$\frac{1}{273}$$


1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

LP Filters: Gaussian

- Effects of sigma:



Original



sigma=1
5 x 5



sigma=2
9 x 9



sigma=4
15 x 15

Types of Noise

- Noise is often modeled as:
 - Gaussian (mean, sigma)
 - Salt-and pepper (0 or 255 with prob p)



Original



Gaussian
sigma=20
mean=0



Salt-and-pepper
 $p=0.3$

Nonlinear Filters

- Linear Filters are good for removing Gaussian noise, but not good for salt-and-pepper (impulsive) noise
- Median filter, alpha-trimmed-mean filter better for impulsive noise

Median, Max, Mean Filters

123	125	126	130	140
122	124	126	127	135
118	120	150	125	134
119	115	119	123	133
111	116	110	120	130

Neighbourhood values:

115, 119, 120, 123, 124,
125, 126, 127, 150

Median value: 124

original

noisy

5 x 5

7 x 7



Alpha-trimmed-mean

- Remove the highest and lowest scores, then take average of the rest

123	125	126	130	140
122	124	126	127	135
118	120	150	125	134
119	115	119	123	133
111	116	110	120	130

Neighbourhood values:

115, 119, 120, 123, 124,
125, 126, 127, 150

Avg: 123.6

Median value: 124