



BMB5113

COMPUTER VISION

DESCRIPTORS

Local Features and Alignment



- To match or align images
 - Global methods sensitive to occlusion, lighting, parallax effects.
 - So look for local features that match well.
 - How would you do it by eye?

Local Features and Alignment

- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align images




Main Questions

- Where will the interest points come from?
 - What are salient features that we'll *detect* in multiple views?
- How to *describe* a local region?
- How to establish *correspondences*, i.e., compute matches?

Local Features and Alignment

- Problem 1:
 - Detect the *same* point *independently* in both images



no chance to match!

We need a **repeatable** detector

Local Features and Alignment

- Problem 2:
 - For each point correctly recognize the corresponding one




We need a reliable and **distinctive** descriptor

Local Descriptors

- There are a lot of methods to detect points.

How to *describe* them for matching?



Point descriptor should be:

1. Invariant
2. Distinctive

Local Descriptors



- Simplest descriptor: list of intensities within a patch.
- What is this going to be invariant to?

Write regions as vectors

$A \rightarrow \mathbf{a}, B \rightarrow \mathbf{b}$


$$\text{NCC} = \frac{\mathbf{a} \cdot \mathbf{b}}{|\mathbf{a}| |\mathbf{b}|}$$

$$-1 \leq \text{NCC} \leq 1$$







Feature Descriptors

- Disadvantage of patches as descriptors:
 - Small shifts can affect matching score a lot



- Solution: histograms



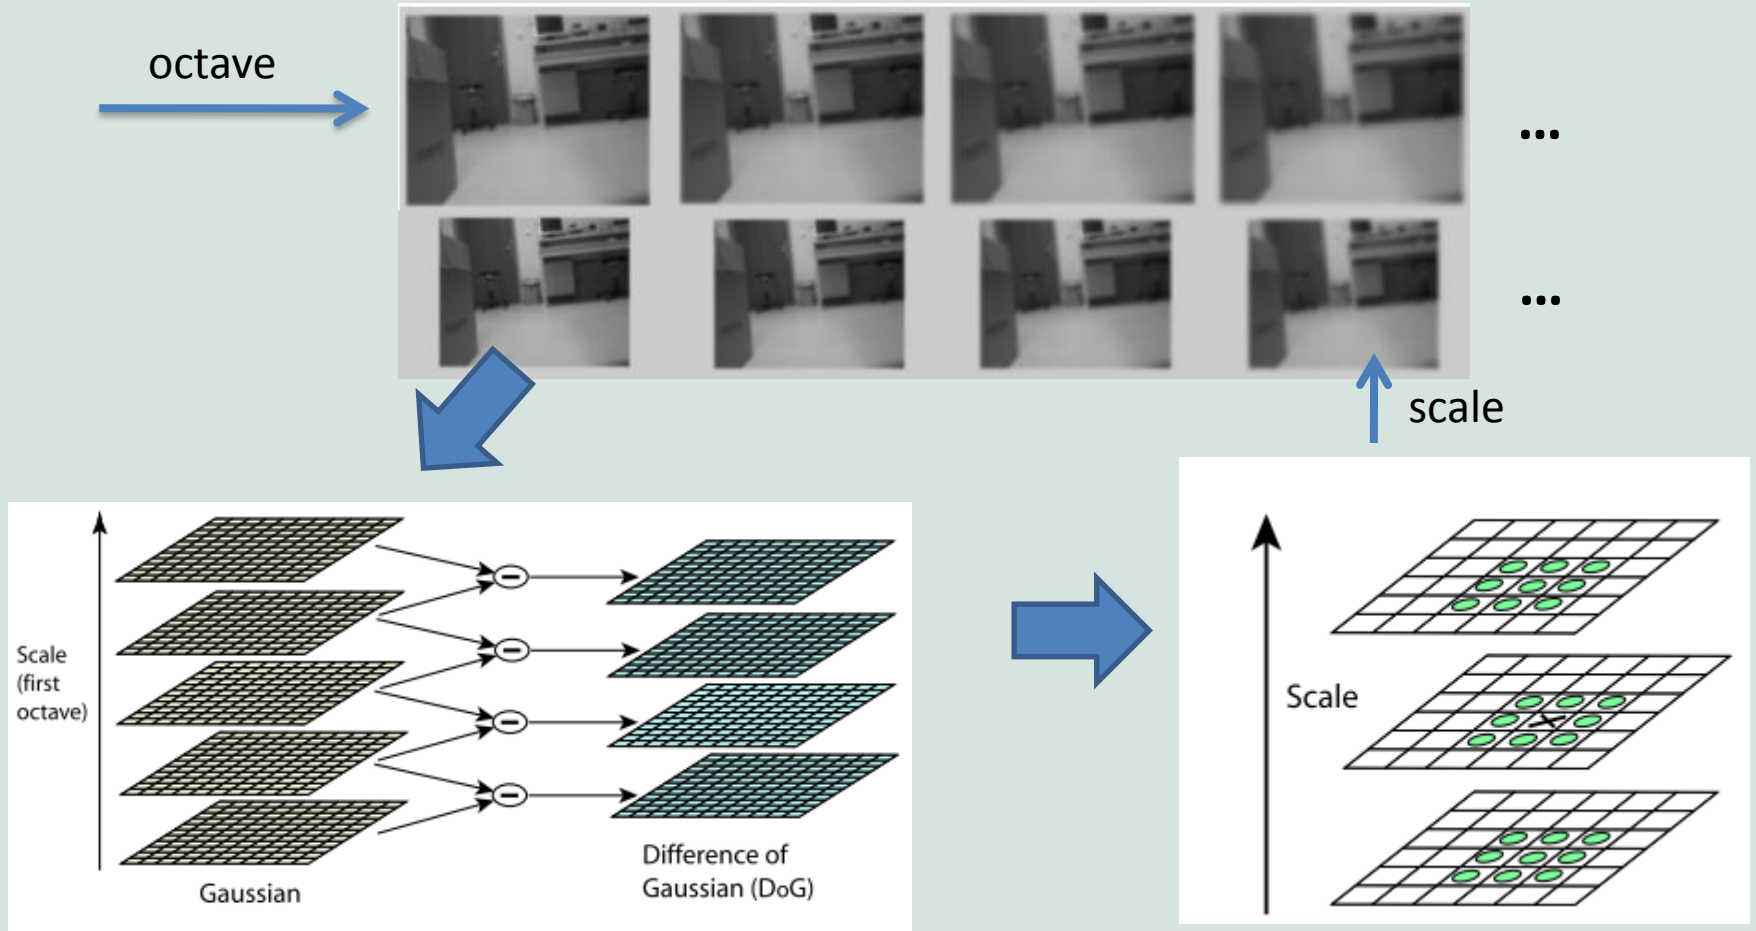
Detectors and Descriptors

- Both detector and descriptor
 - SIFT (Scale Invariant Feature Transform)
 - Detector depends on LoG approximated by Difference of Gaussians (DoG)
 - Descriptor uses histograms of oriented gradients
 - SURF (Speeded-Up Robust Features)
 - Detector depends on determinant of Hessian (DoH)
 - Descriptor uses Haar-wavelets
 - ORB (Oriented FAST and Rotated BRIEF)
- Only descriptor
 - HoG (Histogram of Gradients)
 - depends on statistical information of gradients
 - LBP (Local Binary Patterns)
 - Depends on intensity variations in circular paths

SIFT ALGORITHM

- Extract features invariant to
 - Scale or orientation
 - Affine transformations and illumination (partially)
- Steps involve
 - Scale Space Extrema Detection
 - Keypoint Localization
 - Orientation Assignment
 - Keypoint Descriptor


Scale Space Extrema Detection



- Starting from a σ_0 and using $\sigma_{i+1} = \sqrt{2}\sigma_i$, where $i=0,1,2,3,\dots,7$, compute $I^{\sigma_i} = I * G^{\sigma_i}$
- Halve or double the image size and repeat the above procedure.

Smoothing with a Gaussian

- Parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.




```
fig, ax = plt.subplots(nrows=1, ncols=4)
plt.gray()
for sig in range(1,11,3):
    out = snd.filters.gaussian_filter(im,sig)
    indexAxis = math.floor((sig-1)/3)+1
    plt.subplot(1,4,indexAxis), plt.imshow(out)
    plt.title((r'Gauss blur, $\sigma$='+str(sig)))
plt.show()
```

Effect of σ on Derivatives

- The apparent structures differ depending on Gaussian's scale parameter.
 - Larger values: larger scale edges detected
 - Smaller values: finer features detected




$\sigma = 1$ pixel



$\sigma = 3$ pixels

Laplacian-of-Gaussian (LoG)



$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma^3$$

- Interest points are local maxima or local minima in scale spaces of either
 - Laplacian-of-Gaussian or
 - Difference-of-Gaussian



σ^5

σ^4

σ^3

σ^2

σ




- 26 neighbourhood is checked to determine local maxima or minima

\Rightarrow List of
(x, y, σ)

Keypoint Localization

- Pixels containing extreme values are approximate maxima and minima
- the interpolated (subpixel) locations of the maxima or minima are determined
- a 3D quadratic function is fit around each approximate extremum
- Extreme points can be found differentiating this function and equating it to 0.
 - Taylor series function where D and its derivatives are evaluated at (x, y, σ)
 - $D(x) \rightarrow$ Difference of Gaussian images
 - $\mathbf{x} = (x, y)$ coordinate of extremum
 - Maximum value of this quadratic function is obtained at subpixel location x_s, y_s

$$D(\mathbf{x}) = D + \frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}$$



Keypoint Localization

- A refinement procedure is followed to have more robust interest points
- Determinant of Hessian is used to eliminate edge responses and maintain corners.
 - L_{xx} , L_{xy} , and L_{yy} are second order derivative functions with respect to x or y
 - Refinement condition is checked using trace and determinant of Hessian

$$H(\sigma_i) = \begin{bmatrix} L_{xx}^{\sigma_i} & L_{xy}^{\sigma_i} \\ L_{yx}^{\sigma_i} & L_{yy}^{\sigma_i} \end{bmatrix} \quad 0 \leq \frac{\text{tr}(H(\sigma_i))^2}{\det(H(\sigma_i))} \leq 12$$

Orientation Assignment


- In a local region around the keypoint, gradient magnitudes and directions are computed for each pixel
 - window size is typically 1.5 times the keypoint scale
- Histogram of oriented gradients is generated.
 - 36 bin histogram for better resolution
 - the most prominent one or more orientations in that region are assigned to the keypoint
 - the midpoint of the bin with greatest magnitude is selected

Magnitude:

$$m(x,y) = \sqrt{(L(x+1,y) - L(x-1,y))^2 + (L(x,y+1) - L(x,y-1))^2}$$



Orientation:

$$\tan^{-1} \left(\frac{L(x,y+1) - L(x,y-1)}{L(x+1,y) - L(x-1,y)} \right)$$



Keypoint Descriptor

- A way to distinguish each keypoint from others uniquely
- Again in a window resized according to the keypoint scale, compute the descriptor
 - the image gradient magnitudes and orientations are sampled around the keypoint location
 - the neighbours are Gaussian weighted and window is divided into 4x4 cells
 - Histogram of gradients are computed for each cell using 8 bins and normalized
 - the descriptor coordinates and the gradient orientations are rotated relative to the keypoint orientation



In SIFT:


The window is split into
16 (4×4) cells

Keypoint descriptor length:
16 cells x 8 bins = 128

Rotation Invariant Descriptors

- Find local orientation

Dominant direction of gradient for the image patch



- Rotate patch according to this angle

This puts the patches into a canonical orientation.



Scale Invariant Feature Transform (SIFT)

- Describes image features that are
 - invariant to
 - scale
 - rotation
 - partially invariant to
 - change in illumination
 - 3D camera viewpoint
 - occlusion, clutter
- Fast and efficient—can run close to real time
- Lots of code available
- Provides a basis for object and scene recognition.

Object Recognition with SIFT



A parallelogram -> each recognized object's boundaries
Smaller squares -> the keypoints that were used for recognition.

Scale-Space Blob Detector: Example




Scale-Space Blob Detector: Example




sigma = 9.5859

Scale-Space Blob Detector: Example



Key Point Detection: Example



- (a)** 233x189 image
- (b)** 832 DOG extrema
- (c)** 729 left after peak value threshold
- (d)** 536 left after testing ratio of principle curvatures (removing edge responses)

Working with SIFT Descriptors


- One image yields:
 - n 128-dimensional descriptors: each one is a histogram of the gradient orientations within a patch
 - $[n \times 128 \text{ matrix}]$
 - n scale parameters specifying the size of each patch
 - $[n \times 1 \text{ vector}]$
 - n orientation parameters specifying the angle of the patch
 - $[n \times 1 \text{ vector}]$
 - n 2d points giving positions of the patches
 - $[n \times 2 \text{ matrix}]$
- Remember each patch is related to a keypoint.

Speeded-up Robust Features (SURF)

- a scale and rotation invariant interest point detector and descriptor
- Method:
 - Detector was based on the Hessian matrix but relied on the integral images to reduce the computation time
 - Hessian matrix → a matrix of second order derivatives
 - The descriptor described the distribution of Haar-wavelet responses within the interest point neighborhood
 - Generally only 64 dimensions were used reducing the time for feature computation and modeling.

Speeded-up Robust Features (SURF)

- Results:
 - sometimes more than 10% improvement in recall for the same level of precision
 - For the same number of points, computations of the detector and the descriptor
 - 391 ms for SURF-128
 - 1036 ms for SIFT
- approximated or outperformed previously proposed approaches respect to
 - repeatability
 - distinctiveness and robustness
 - computation and comparison time




Feature Matching


- Given a feature in I_1 , how to find the best match in I_2 ?
 1. Define distance function that compares two descriptors
 2. Test all the features in I_2 , find the one with min distance

Feature Distance

- How to define the difference between two features f_1, f_2 ?
 - Simple approach is $\text{SSD}(f_1, f_2)$
 - sum of square differences between entries of the two descriptors
 - can give good scores to very ambiguous (bad) matches




I_1




I_2

Feature Distance

- How to define the difference between two features f_1, f_2 ?
 - Better approach: ratio distance = $\text{SSD}(f_1, f_2) / \text{SSD}(f_1, f'_2)$
 - f_2 is best SSD match to f_1 in I_2
 - f'_2 is 2nd best SSD match to f_1 in I_2
 - gives large values for ambiguous matches



I_1



I_2

Matching Keypoints

SIFT



SURF



Feature Detector and Descriptors

Summary

- Stable (repeatable) feature points can be detected regardless of image changes
 - Scale: search for correct scale as *maximum* of appropriate function
 - Affine: approximate regions with *ellipses* (this operation is affine invariant)
- Invariant and distinctive descriptors can be computed
 - Using invariant *moments*
 - *Normalizing* with respect to scale and affine transformation

Evaluating the Matching Results

- How can we measure the performance of a feature matcher?
 - True/false positives
- The distance threshold affects performance
 - True positives = # of detected matches that are correct
 - Suppose we want to maximize these—how to choose threshold?
 - False positives = # of detected matches that are incorrect
 - Suppose we want to minimize these—how to choose threshold?

Evaluating the Matching Results

- How can we measure the performance of a feature matcher?


$$\frac{\text{\# true positives}}{\text{\# matching features (positives)}}$$

$$\text{True positive rate (recall)} = \frac{TP}{TP+FN}$$

$$\text{False positive rate} = \frac{FP}{FP+TN}$$

$$\text{Precision} = \frac{TP}{TP+FP}$$

$$\text{F1-score} = \frac{2 * \text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{2TP}{2TP+FP+FN}$$




Evaluating the Matching Results

- How can we measure the performance of a feature matcher?

ROC curve

« Receiver Operator Characteristic »

$$\frac{\text{\# true positives}}{\text{\# matching features (positives)}}$$



- ROC Curves
 - Generated by counting # correct matches / # incorrect matches, for different thresholds
 - Want to maximize area under the curve (AUC)
 - Useful for comparing different feature matching methods


Histogram of Oriented Gradients (HOG)

- Compress image to 64x128 pixels
- Convolution with $[-1 \ 0 \ 1]$ $[-1;0; 1]$ filters
- Compute gradient magnitude + direction
- For each pixel
 - take the color channel with greatest magnitude as final gradient





HoG: Cell Histograms

- Divide the image to cells, each cell 8x8 pixels
- Snap each pixel's direction to one of 18 gradient orientations
- Build histogram per-cell using magnitudes





HoG: Histogram Interpolation


- Interpolated trilinearly:
 - Bilinearly into spatial cells
 - Linearly into orientation bins



HoG: Normalization



Current cell : 1x18 histogram



Block: 2x2 cell
overlapping with current cell

1. **contrast sensitive features:** 18 orientation -> 18 dim
2. **contrast insensitive features:** 9 orientation -> 9 dim


Normalize 4 times by its neighbor blocks, and average them

3. **texture features:** sum of the magnitude over all orientation and normalize 4 times by its neighbor blocks, not average -> 4 dim




In total each cell : 18+9+4 dimensions of feature

Final Descriptor

- Concatenation the normalized histogram




Visualization:



HOG Descriptor:

1. **Compute gradients** on an image region of 64×128 pixels
2. **Compute histograms** on «cells» of typically 8×8 pixels (i.e. 8×16 cells)
3. **Normalize histograms** within overlapping blocks of cells
4. **Concatenate histograms**



It is a typical procedure of feature extraction !

Local Binary Patterns

- A non-parametric feature extraction method (Ojala v.d., 1994, 1996)
- Analyzes texture pattern and gives orthogonal measures of local contrast
- LBP is
 - Highly tolerant to different illumination conditions
 - Sensitive to small changes in gray-level image
 - Low cost to compute

LBP Algorithm


- Input image
 1. Divide image into cells (e.g. every cell has 16×16 pixels)
 2. Every pixel in a cell is compared to its neighbors
 - Neighborhood can be in (8, R) or (16, R) structure where R is distance to neighbors and 8 or 16 is the number of neighbours.
 - Follow pixels along a circular path in clockwise or counter clockwise direction
 - If central pixel is greater than or equal to neighbor, label as 0 otherwise label as 1
 - As a result for 8 neighborhood an 8-digit binary number is obtained
 - This binary number is generally converted to a decimal number
 3. A frequency histogram of the decimal numbers is computed in each cell
 - A 256-length feature histogram is obtained
 4. 256-length feature histogram is converted to 59-length uniform pattern
 5. Normalized histograms of each cells is combined to have a feature vector for the whole image

LBP Neighborhood

187	177	166		1	0	0		
188	178	165	→	1		0	→	(10000111) =135
188	182	168		1	1	0		

$$\text{LBP}_{P,R}(x_c) = \sum_{p=0}^{P-1} u(x_p - x_c) 2^p, \quad u(y) = \begin{cases} 0, & \text{if } y < 0 \\ 1, & \text{if } y \geq 0 \end{cases}$$


- Circular (8, 1), (16, 2) and (8, 2) neighborhoods



LBP Patterns

- For a P-bit binary number there are 2^P different LBP codes
 - 8 bit number $\rightarrow 2^8=256$ different patterns
 - LBP patterns can be
 - uniform
 - non-uniform
- Uniform LBP patterns
 - An extension of LBP
 - more robust to noise and computationally less complex
 - yield better detection or recognition results
 - also related to rotational invariance issue
 - If number of passes from «0» to «1» less than or equal to two
 - 0000000 (0 pass) \rightarrow uniform
 - 00011100 (2 passes) \rightarrow uniform
 - 11101111 (2 passes)) \rightarrow uniform
 - 00110100 (4 passes) \rightarrow nonuniform
 - 01010100 (6 passes)) \rightarrow nonuniform

58 Different Uniform Patterns for (8, R) Neighborhood




- Every uniform pattern has a different label
- Every nonuniform pattern assigned to a single label
- For 8-bit pattern
 - 59 labels

(Pietikäinen et al., 2011)

Some Structures Detected by LBP

- Black → bit 1
- White → bit 0
- Gray → central pixel




Oriented FAST and Rotated BRIEF (ORB)

- Combines FAST feature detectors and BRIEF descriptors, (Ethan Rublee, 2011)
- Enables fast detections and descriptions of features
 - Useful for real-time applications such as robotics vision systems or smartphone applications

Features from Accelerated Segment Test (FAST)

- Uses pixel neighborhood to compute key points in an image
- Algorithm
 - An interesting point candidate pixel (i,j) is selected with an intensity $I(i,j)$
 - In a circle of 16 pixels, given a threshold t ,
 - estimate n adjoining points which are brighter or darker than pixel (i,j) by the threshold t .
 - n is chosen as 12
 - In a high-speed test only four pixels at 1, 9, 5, 13 in the figure are checked. Intensity values of at least three of these pixels determine whether (i,j) is a corner.



FAST Features with Orientation

- FAST does not produce rotation information
- ORB uses FAST features with orientation information
- Using a circular radius of 9 pixels a vector between the computed intensity centroid and center of the corner is used to describe orientation
- In circular region the moments are computed

$$m_{p,q} = \sum_{x,y} x^p y^q I(x,y)$$

- Using computed moments intensity centroid is determined

$$C = \left(\frac{m_{1,0}}{m_{0,0}}, \frac{m_{0,1}}{m_{0,0}} \right)$$

- Patch center O and intensity centroid C are joined to form the orientation vector \overrightarrow{OC} .
- The orientation of the patch then simply is:

$$\theta = \text{atan2}(m_{0,1}, m_{1,0})$$

Binary Robust Independent Elementary Features (BRIEF)

- A descriptor proposed by Michael Calonder et al. (2010)
- Unlike relatively high dimensional features of
 - SIFT: 128 x 4 byte
 - SURF: 64 x 4 byte
 - BRIEF consumes less memory
 - 128, 256 or 512 bits
 - 16, 32, 64 bytes
- Similar to LBP, BRIEF computes intensity differences in a small patch of image and represents them as a binary string

Lots of Applications

- Features are used for:
 - Image alignment (e.g., mosaics)
 - 3D reconstruction
 - Motion tracking
 - Object recognition
 - Indexing and database retrieval
 - Robot navigation
 - ... other