



BMB5113

COMPUTER VISION




TRANSFORMATIONS & ALIGNMENT

Stitching Problem




Matching Image Features

- 1) **Feature Detection:**
Identify image features
- 2) **Feature Description:**
Extract feature descriptor for each feature
- 3) **Feature Matching:**
Find candidate matches between features
- 4) **Feature Correspondence:**
Find consistent set of (inlier) correspondences between features



An Example




Problem 1: Preprocessing

- Most feature descriptors only work with grayscale images

RGB image: 584 x 778 x 3

Grayscale image: 584 x 778

- Convert color images to grayscale



Problem 2: Detecting Keypoints


- Identify features in an image

Detector: Hessian, Harris, FAST, etc...

Independently detect keypoints in both images


Problem 3: Extracting Descriptors

- Extract feature descriptors for each keypoint
 - Derived from pixels surrounding an interesting point.
 - HoG, LBP, Haar-wavelets, etc...
 - Remember SIFT, SURF, ORB



Problem 4: Matching Features


- Find pairs of features in the two images that match
 - Possible strategies: Match Threshold, Nearest Neighbor Symmetric, and Nearest Neighbor Ratio
- Find pairs of indices in feature list of image 1 and feature list of image 2 that match



Problem 5: RANSAC to Estimate Homography

- Exclude outlier matches
- Compute a homography to map one image plane to the other
 - Solution is RANSAC (Random Sample Consensus)
- RANSAC
 - Finds inlier points and compute a transformation from image 2 to image 1
 - Returns the transformation from inlier_points2 to inlier_points1
 - Transform_type?
 - Similarity, Affine, Projective

Problem 5: RANSAC to Estimate Homography



Problem 6: Stitching Panorama

- Warp the two images to make a final panoramic image



Alignment

- Homographies
- Rotational Panoramas
- RANSAC
- Global alignment
- Warping
- Blending



(a)




(b)





(c)

Motivation: Recognition



Motivation: Medical Image Registration




Motivation: Mosaics

- Getting the whole picture
 - Consumer camera: $50^\circ \times 35^\circ$




Motivation: Mosaics

- Getting the whole picture
 - Consumer camera: $50^\circ \times 35^\circ$
 - Human Vision: $176^\circ \times 135^\circ$



Motivation: Mosaics

- Getting the whole picture
 - Consumer camera: $50^\circ \times 35^\circ$
 - Human Vision: $176^\circ \times 135^\circ$



- Panoramic Mosaic = up to $360^\circ \times 180^\circ$

Motion models


- What happens when we take two images with a camera and try to align them?
- translation?
- rotation?
- scale?
- affine?
- perspective?



Image Warping

- image filtering: change *range* of image

- $g(x) = h(f(x))$



- image warping: change *domain* of image

- $g(x) = f(h(x))$





Image Warping


- image filtering: change *range* of image

- $g(x) = h(f(x))$



- image warping: change *domain* of image

- $g(x) = f(h(x))$



Parametric (Global) Warping

- Examples of parametric warps:



translation



rotation



aspect



affine




perspective



cylindrical


Image Warping

- Given a coordinate transform $\mathbf{x}' = \mathbf{h}(\mathbf{x})$ and a source image $\mathbf{f}(\mathbf{x})$, how do we compute a transformed image $\mathbf{g}(\mathbf{x}') = \mathbf{f}(\mathbf{h}(\mathbf{x}))$?




Forward Warping

- Send each pixel $f(x)$ to its corresponding location $x' = h(x)$ in $g(x')$
 - What if pixel lands “between” two pixels?




Forward Warping

- Send each pixel $f(x)$ to its corresponding location $x' = h(x)$ in $g(x')$
 - What if pixel lands “between” two pixels?
 - Solution: add “contribution” to several pixels, normalize later (splatting)




Inverse Warping

- Get each pixel $g(x')$ from its corresponding location $x' = h(x)$ in $f(x)$
 - What if pixel comes from “between” two pixels?



Inverse Warping



Get each pixel $g(x', y')$ from its corresponding location
 $(x, y) = h^{-1}(x', y')$ in the first image

Q: what if pixel comes from “between” two pixels?

A: *Interpolate* color value from neighbors

– nearest neighbor, bilinear...


Interpolation

- Possible interpolation filters:
 - nearest neighbor
 - bilinear
 - bicubic (interpolating)
 - sinc / FIR
- Needed to prevent “jaggies” and “texture crawl”



Bilinear interpolation

- Sampling at $f(x, y)$:



$$\begin{aligned} f(x, y) = & (1 - a)(1 - b) f[i, j] \\ & + a(1 - b) f[i + 1, j] \\ & + ab f[i + 1, j + 1] \\ & + (1 - a)b f[i, j + 1] \end{aligned}$$

Python Warp Function

```
import cv2

translation_matrix = np.float32([[1,0,160],[0,1,40]])

transformed = cv2.warpAffine(img, translation_matrix,
                              (img.shape[1], img.shape[0]))

affine_tr = cv2.getAffineTransform(pts1,pts2)
transformed = cv2.warpAffine(img, affine_tr, (shape1,
shape2))

from skimage import transform as tf

img2 = tf.rotate(img1, 180)

tform = tf.AffineTransform(scale=(1.3, 1.1),
rotation=0.5,translation=(0, -200))


img3 = tf.warp(img1, tform)
```

Transformations

- How to make and use correspondences between
 - images and maps,
 - images and models,
 - images and other images
- Transformation may be in two or three dimensions

Image Registration

- Points of two images with similar viewpoints of essentially the same scene are geometrically transformed so that
 - corresponding feature points of the two images have the same coordinates after transformation



$$M[x, y] \cong I[g(x, y), h(x, y)]$$

$$I[r, c] \cong M[g^{-1}(r, c), h^{-1}(r, c)]$$

Representation of Points

- A 2D point has two coordinates and is conveniently represented as either
 - a row vector $P=[x,y]$
 - column vector $P=[x,y]^t$
- Homogeneous Coordinates:
 - A convenient notation for computer processing of points
 - especially when affine transformations are used.
- The homogeneous coordinates of a 2D point $P=[x,y]^t$ are $[s.x,s.y,s]^t$, where s is a scale factor and commonly 1.0

2D Coordinate Transformations

- translation: $\mathbf{x}' = \mathbf{x} + \mathbf{t}$ $\mathbf{x} = (x, y)$
- rotation: $\mathbf{x}' = \mathbf{R} \mathbf{x} + \mathbf{t}$
- similarity: $\mathbf{x}' = s \mathbf{R} \mathbf{x} + \mathbf{t}$
- affine: $\mathbf{x}' = \mathbf{A} \mathbf{x} + \mathbf{t}$
- perspective: $\underline{\mathbf{x}}' \cong \mathbf{H} \underline{\mathbf{x}}$ $\underline{\mathbf{x}} = (x, y, 1)$
($\underline{\mathbf{x}}$ is a *homogeneous* coordinate)
 - These all form a nested *group* (closed with inverse)

Basic 2D Transformations

- Basic 2D transformations as 3x3 matrices

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta & 0 \\ \sin \Theta & \cos \Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$


Rotate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Shear

Scaling

- Uniform scaling
 - changes all coordinates in the same way
 - equivalently changes the size of all objects in the same way.
- A 2D point $P=[1,2]$ scaled by a factor of 2 to obtain the new point $P'=[2,4]$.




Scaling

- Scaling is a linear transformation
 - can be easily represented in terms of the scale factor applied to the two basis vectors for 2D Euclidean space.
- For example,
 $[1,2]=1[1,0]+2[0,1]$ and
 $2[1,2]=2(1[1,0]+2[0,1])=[2,4]$

Rotation

- A second common operation is rotation about a point in 2D space.
- A 2D point $P=[x,y]$ rotated by angle θ counterclockwise about the origin to obtain the new point $P'=[x',y']$.



Rotation

- Rotation is also a linear transformation,
 - the columns of the matrix are result of the transformation applied to the basis vectors
 - transformation of any other vectors can be expressed as a linear combination of the basis vectors.

$$\begin{aligned} R_{\theta}(x \cdot [1,0] + y \cdot [0,1]) &= (x \cdot R_{\theta} \cdot [1,0]) + (y \cdot R_{\theta} \cdot [0,1]) \\ &= x[\cos(\theta), \sin(\theta)] + y[-\sin(\theta), \cos(\theta)] \end{aligned}$$

$$[x', y'] = [x \cos(\theta) - y \sin(\theta), x \sin(\theta) + y \cos(\theta)]$$

$$[x', y'] = \begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} x \cos(\theta) - y \sin(\theta) \\ x \sin(\theta) + y \cos(\theta) \end{bmatrix}$$

Orthogonal and Orthonormal Transformations

- **Orthogonal** A set of vectors is orthogonal if all pairs of vectors in the set are perpendicular (have scalar product of zero).
- **Orthonormal** A set of vectors is orthonormal if it is an orthogonal set and all vectors have unit length.
- *A rotation preserves both the length of the basis vectors and their orthogonality.*

Translation

- Point coordinates need to be shifted by some constant amount.
 - equivalent to changing the origin of the coordinate system.
- Since translation does not map the origin $[0,0]$ to itself, we cannot model it using a simple 2×2 matrix as has been done for scaling and rotation.
 - In other words, **translation is not a linear operation.**

Translation

- Translation can be represented as the matrix multiplication in homogeneous coordinates.
- A translation T of point $[x,y]$ so that $[x',y'] = T([x,y]) = [x+t_x, y+t_y]$

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + t_x \\ y + t_y \\ 1 \end{bmatrix}$$

2D Affine Transformations


$$\begin{bmatrix} x' \\ y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- Affine transformations are combinations of ...
 - Linear transformations (rotation + scale), and
 - Translations
- Properties of affine transformations:
 - Origin does not necessarily map to origin
 - Lines map to lines
 - **Parallel lines remain parallel**
 - Ratios are preserved
 - Closed under composition
 - Models change of basis

Projective Transformations

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- Projective transformations:
 - Affine transformations, and
 - Projective warps
- Parallel lines do not necessarily remain parallel




Fitting an Affine Transformation



Affine model approximates perspective projection of planar objects.

Points in 3D Space

- A 3D point (x,y,z) – x , y , and z coordinates
- Column vectors are used to represent points
 - Homogeneous coordinates of a 3D point $(x,y,z,1)^T$
- In general homogeneous coordinates in 3D
 - $[x,y,z,1]^T \equiv (x,y,z,w)$
- Transformation will be performed using 4x4 matrix



$$\begin{bmatrix} n_x & o_x & a_x & p_x \\ n_y & o_y & a_y & p_y \\ n_z & o_z & a_z & p_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} n_x \cdot x + o_x y + a_x z + p_x \\ n_y \cdot x + o_y y + a_y z + p_y \\ n_z \cdot x + o_z y + a_z z + p_z \\ 1 \end{bmatrix}$$

Affine Transformations in 3-D

- Translation

$$T(d, h, l) = \begin{bmatrix} 1 & 0 & 0 & d \\ 0 & 1 & 0 & h \\ 0 & 0 & 1 & l \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x + d \\ y + h \\ z + l \\ 1 \end{bmatrix}$$

- Scaling


$$S(a, f, k) = \begin{bmatrix} a & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & k & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} ax \\ fy \\ kz \\ 1 \end{bmatrix}$$

- Rotation?

$$\begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x \cos \theta - y \sin \theta \\ x \sin \theta + y \cos \theta \\ 1 \end{bmatrix}$$

3D Rotation


- 3D rotation is done around an axis
- Fundamental rotations – rotate about x, y, or z axes
- Counter-clockwise rotation is referred to as positive rotation
 - when you look from positive down to negative axis



Rotation about Z axis

- Same with 2-D version except the added additional row and column for the z axis.

$$\begin{bmatrix} \cos \theta & -\sin \theta & 0 & 0 \\ \sin \theta & \cos \theta & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \cos \theta - y \sin \theta \\ x \sin \theta + y \cos \theta \\ z \\ 1 \end{bmatrix}$$



3D Transformation

- Rotation about y

$$(z \rightarrow y, y \rightarrow x, x \rightarrow z)$$


$$z' = -x \sin(\theta) + z \cos(\theta)$$

$$x' = x \cos(\theta) + z \sin(\theta)$$

$$y' = y$$

$$\begin{vmatrix} \cos(\theta) & 0 & \sin(\theta) & 0 \\ 0 & 1 & 0 & 0 \\ -\sin(\theta) & 0 & \cos(\theta) & 0 \\ 0 & 0 & 0 & 1 \end{vmatrix}$$

$$\begin{bmatrix} x \cos \theta - y \sin \theta \\ x \sin \theta + y \cos \theta \\ z \\ 1 \end{bmatrix}$$



3D Transformation

- Rotation about x

$$(z \rightarrow y, y \rightarrow z, x \rightarrow x)$$



$$y' = y \cos(\theta) - z \sin(\theta)$$

$$z' = y \sin(\theta) + z \cos(\theta)$$

$$x' = x$$

$$\begin{vmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos(\theta) & -\sin(\theta) & 0 \\ 0 & \sin(\theta) & \cos(\theta) & 0 \\ 0 & 0 & 0 & 1 \end{vmatrix}$$

$$\begin{bmatrix} x \cos \theta - y \sin \theta \\ x \sin \theta + y \cos \theta \\ z \\ 1 \end{bmatrix}$$



The 3 Rotation Matrices

- Rotation about an arbitrary axis can be represented by only these three main matrices.

$$R_z(\theta) = \begin{bmatrix} \cos \theta & -\sin \theta & 0 & 0 \\ \sin \theta & \cos \theta & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$R_x(\theta) = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta & 0 \\ 0 & \sin \theta & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$R_y(\theta) = \begin{bmatrix} \cos \theta & 0 & \sin \theta & 0 \\ 0 & 1 & 0 & 0 \\ -\sin \theta & 0 & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Skew or Shear in 3-D

- In shearing x, y or z coordinates can be updated depending on the values of other point coordinates.

$$H_x(b) = \begin{bmatrix} 1 & b & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Matrix for Perspective Projection?

- Division is needed to do projection!
 - But, matrix multiplication only does multiplication and addition
- What about:

$$M_{\text{per}} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 0 \end{bmatrix}$$


$$M_{\text{per}} \cdot P = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 0 \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ z \\ z/d \end{bmatrix}$$

i-Point Perspective Transformation





1 - point

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 1 \end{bmatrix}$$




2 - point

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & p & q & 1 \end{bmatrix}$$




3 - point

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ p & q & r & 1 \end{bmatrix}$$



Fitting an Affine Transformation

- Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

Fitting an Affine Transformation

- Assuming we know the correspondences, how do we get the transformation?

$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$


$$\begin{bmatrix} \vdots \\ \vdots \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \vdots \\ \vdots \end{bmatrix}$$

Fitting an Affine Transformation

$$\begin{bmatrix} \dots & & & & & \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$

- How many matches (correspondence pairs) do we need to solve for the transformation parameters?
- Once we have solved for the parameters, how do we compute the coordinates of the corresponding point for
? (x_{new}, y_{new})

Panoramas



...




image from S. Seitz

Obtain a wider angle view by combining multiple images.

How to Stitch Together a Panorama?

- Basic Procedure
 - Take a sequence of images from the same position
 - Rotate the camera about its optical center
 - Compute transformation between second image and first
 - Transform the second image to overlap with the first
 - Blend the two together to form a mosaic
 - If there are more images, repeat
- ...but **wait**, why should this work at all?
 - What about the 3D geometry of the scene?
 - Why aren't we using it?


Panoramas: Generating Synthetic Views



Can generate any synthetic camera view
as long as it has **the same center of projection!**

Image reprojection

- The mosaic has a natural interpretation in 3D
 - The images are reprojected onto a common plane
 - The mosaic is formed on this plane
 - Mosaic is a *synthetic wide-angle camera*




Homography

- How to relate two images from the same camera center?
 - how to map a pixel from PP1 to PP2?
- Think of it as a 2D **image warp** from one image to another.
- A projective transform is a mapping between any two PPs with the same center of projection
 - rectangle should map to arbitrary quadrilateral
 - parallel lines aren't parallel
 - but must preserve straight lines
 - called **Homography**

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$\mathbf{p}' \qquad \mathbf{H} \qquad \mathbf{p}$



Homography

(x, y)




$$\begin{pmatrix} wx'/w & wy'/w \end{pmatrix} \\ = (x', y')$$

To **apply** a given homography **H**

- Compute $\mathbf{p}' = \mathbf{H}\mathbf{p}$ (regular matrix multiply)
- Convert \mathbf{p}' from homogeneous to image coordinates

$$\begin{bmatrix} wx' \\ wy' \\ w \\ \mathbf{p}' \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ \mathbf{H} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \\ \mathbf{p} \end{bmatrix}$$

Homography



To **compute** the homography given pairs of corresponding points in the images, we need to set up an equation where the parameters of **H** are the unknowns...

Solving for Homographies

$$\mathbf{p}' = \mathbf{H}\mathbf{p}$$

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- Can set scale factor $w=1$. So, there are 8 unknowns.
- Set up a system of linear equations:

$$\mathbf{A}\mathbf{h} = \mathbf{b}$$

where vector of unknowns $\mathbf{h} = [a, b, c, d, e, f, g, h]^T$

- Need at least 8 eqs, but the more the better...
- Solve for \mathbf{h} . If overconstrained, solve using least-squares:

$$\min \|\mathbf{A}\mathbf{h} - \mathbf{b}\|^2$$

- Practical solutions to equation are through functions such as
 - `lmdivide`, `svd`

Image Warping with Homographies






Image Rectification



Analysing Patterns and Shapes




**Automatic
rectification**




**From Martin Kemp, *The Science of Art*
(*manual reconstruction*)**

Changing Camera Center

- Does it still work?




Planar Scene (Far-away Scene)



- PP3 is a projection plane of both centers of projection, so we are OK!
- This is how big aerial photographs are made.




Outliers

- **Outliers** can hurt the quality of our parameter estimates, e.g.,
 - an erroneous pair of matching points from two images
 - an edge point that is noise, or doesn't belong to the line we are fitting.



Least Squares Line Fitting

- Assuming all the points that belong to a particular line are known
- Outliers affect least squares fit



Random Sample Consensus (RANSAC)

- Approach:
 - avoid the impact of outliers
 - look for inliers and use those only
- Intuition:
 - if an outlier is chosen to compute the current fit, then the resulting line won't have much support from rest of the points.

RANSAC


- RANSAC algorithm:
 1. Randomly select a *seed group* of points on which to base transformation estimate (e.g., a group of matches)
 2. Compute transformation \mathbf{H} from seed group
 3. Find *inliers* to this transformation
$$p'_i = H.p_i, p_i \rightarrow {}^m p_i,$$
$$\text{if } \|{}^m p_i - p'_i\| < \text{thresh} \text{ then } (p_i, {}^m p_i) \text{ is inlier to } \mathbf{H}$$

even if $(p_i, {}^m p_i)$ matching pair is incorrect.
 4. If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers
 - Keep the transformation with the largest number of inliers


RANSAC Line Fitting Example




RANSAC Line Fitting Example



RANSAC Line Fitting Example




RANSAC Line Fitting Example



RANSAC Line Fitting Example




RANSAC Line Fitting Example



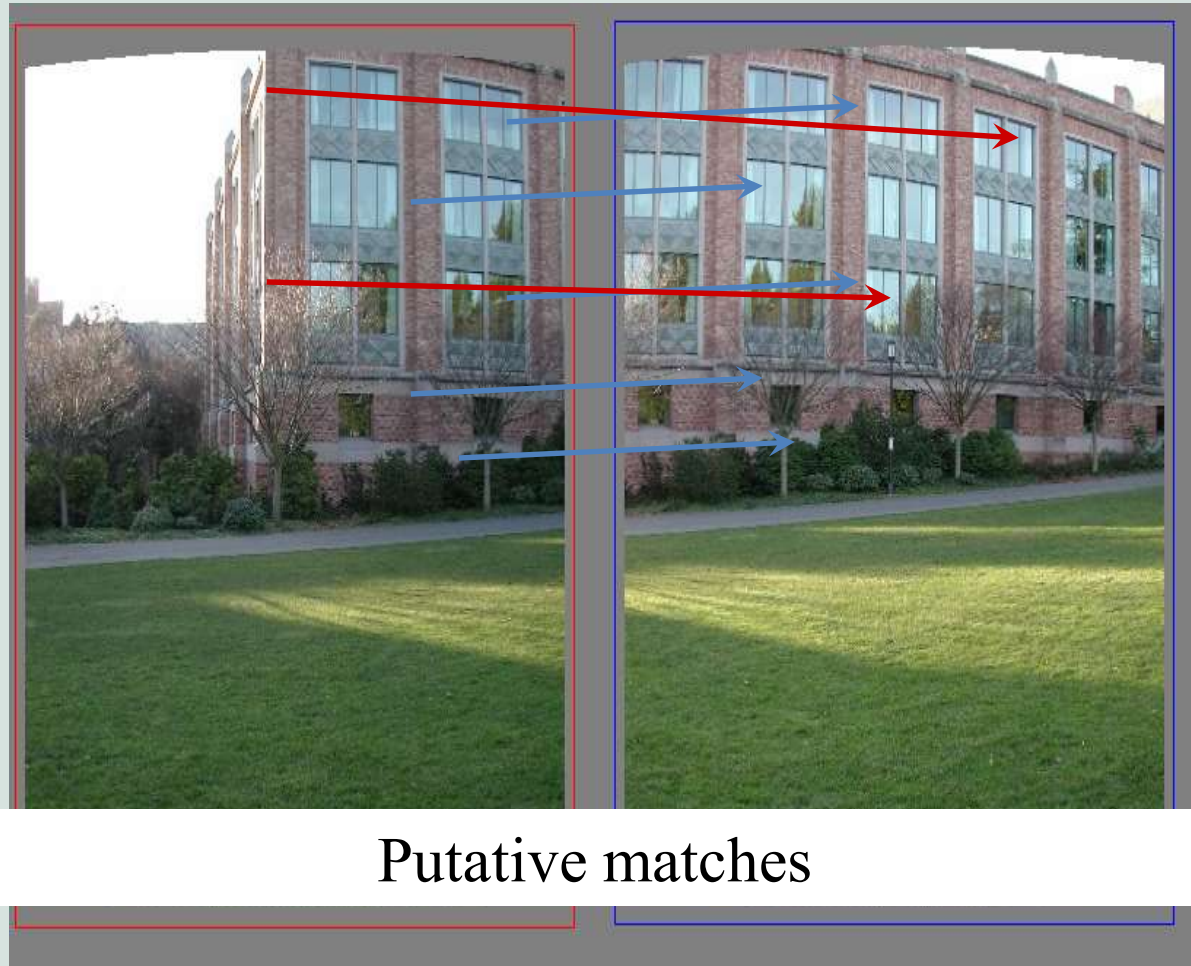
Repeat, until get a
good result

RANSAC Line Fitting Example




Repeat, until get a
good result


RANSAC Example: Translation



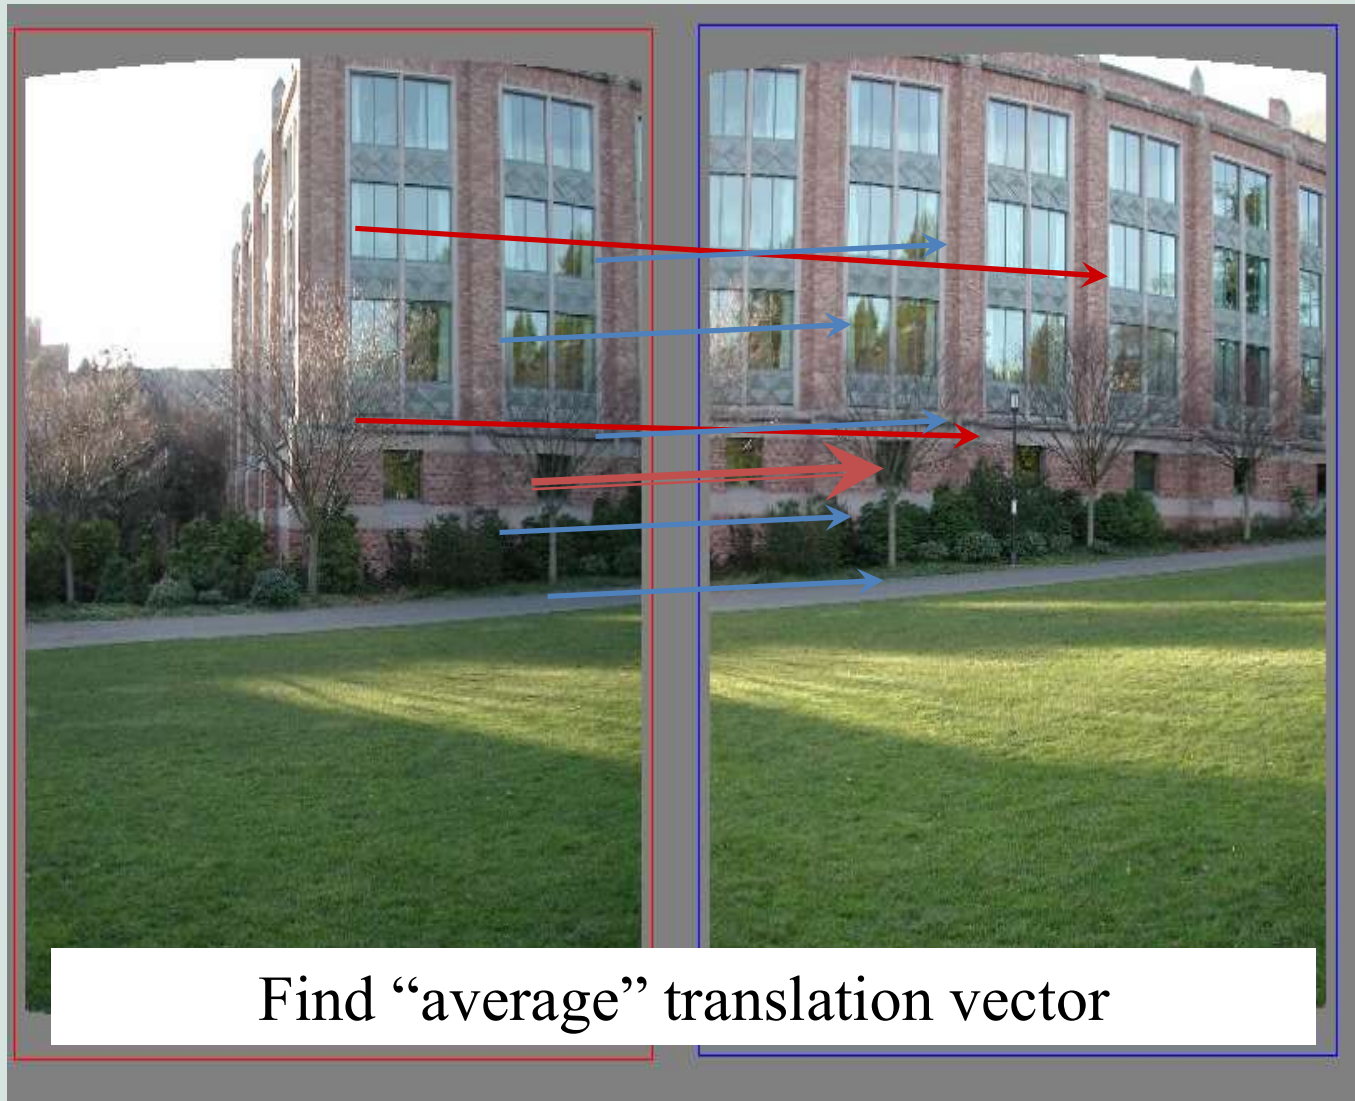
RANSAC Example: Translation



RANSAC Example: Translation



RANSAC example: Translation




Exercise 1: Rotation About A Point

- Give the 3x3 matrix that represents a $\pi/2$ rotation of the plane about the point $[5,8]$.
- Hint:
 - First derive the matrix $D_{-5,-8}$ that translates the point $[5,8]$ to the origin of a new coordinate frame.
 - The matrix which we want will be the combination
$$D_{5,8} R_{\pi/2} D_{-5,-8}$$
- Check that your matrix correctly transforms points $[5,8]$, $[6,8]$ and $[5,9]$

Exercise 2: Rotation, Scaling and Translation

- Given a planar workspace $W[x,y]$
 - An image $I[r,c]$ is taken by a square-pixel camera looking perpendicularly down on a planar workspace $W[x,y]$.
 - This conversion can be done by composing a rotation \mathbf{R} , a scaling \mathbf{S} and a translation \mathbf{D} as given and denoted by ${}^w\mathbf{P}_j = \mathbf{D}_{x_0,y_0} \mathbf{S}_s \mathbf{R}_\theta {}^i\mathbf{P}_j$




Exercise 2: Rotation, Scaling and Translation

- ${}^w\mathbf{P}_j = \mathbf{D}_{x_0, y_0} \mathbf{S}_s \mathbf{R}_\theta {}^i\mathbf{P}_j$

$$\begin{bmatrix} x_w \\ y_w \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & x_0 \\ 0 & 1 & y_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} s & 0 & 0 \\ 0 & s & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

$$x_w = x_i s \cos \theta - y_i s \sin \theta + x_0$$

$$y_w = x_i s \sin \theta + y_i s \cos \theta + y_0$$



Exercise 2:

Rotation, Scaling and Translation

- Control points are clearly distinguishable and easily measured points to establish known correspondences between different coordinate spaces.
- Given the following two matches between the world and image planes

$${}^iP_1 = [100, 60] \text{ and } {}^wP_1 = [200, 100]$$

$${}^iP_2 = [380, 120] \text{ and } {}^wP_2 = [300, 200]$$

- θ is easily determined independent of the other parameters as follows:

$$\theta_i = \arctan(({}^iy_2 - {}^iy_1)/({}^ix_2 - {}^ix_1)) \quad \theta_i = \arctan(60/280) = 12.09^\circ$$

$$\theta_w = \arctan(({}^wy_2 - {}^wy_1)/({}^wx_2 - {}^wx_1)) \quad \theta_w = \arctan(100/100) = 45^\circ$$

$$\theta = \theta_w - \theta_i \quad \theta = 32.91^\circ$$

- Once θ is determined, all *sin* and *cos* elements are known
- There are 4 equations and 3 unknowns which can be solved for s and x_0, y_0 .

$$200 = 100.s.0.84 - 60.s.0.54 + x_0 = 51.6.s + x_0$$

$$100 = 100.s.0.54 - 60.s.0.84 + y_0 = 3.6.s + y_0$$

$$300 = 380.s.0.84 - 120.s.0.54 + x_0 = 254.4.s + x_0$$

$$200 = 380.s.0.54 - 120.s.0.84 + y_0 = 104.4.s + y_0$$